

Natural Language Processing Concepts

Justine Bradley Tulio
Bachelor of Science in Information
Technology
Jose Rizal University
Manila, Philippines
justinebradley.tulio@my.jru.edu

Abstract—This reflection report summarizes my learning experience throughout Unit 5 on Natural Language Processing (NLP) and the accompanying text analytics exercises. Before completing this unit, I understood NLP only at a general level and viewed it mainly as a technology that allows computers to process human language. Through the readings and practical activities, I developed a deeper appreciation of the complete NLP pipeline, from preprocessing textual data to representing language through statistical and semantic methods. More importantly, I realized that each technique serves a specific purpose, and selecting the appropriate approach depends on the problem being solved. This experience not only strengthened my technical understanding but also showed me how these foundational concepts will support more advanced topics in machine learning, large language models, and future software development projects.

Keywords—*Natural Language Processing, Reflection, Text Analytics, Learning Experience, Artificial Intelligence*

I. INTRODUCTION

my understanding of Natural Language Processing was fairly limited. I knew that technologies such as chatbots, virtual assistants, and translation systems relied on NLP, but I had not considered the amount of processing required before a computer could understand human language. I assumed that modern AI models could simply "read" text in much the same way people do. As I progressed through the readings and exercises, I realized that this assumption overlooked many important steps that make language processing possible.

One of the things I appreciated most about this unit was how it gradually introduced the concepts rather than presenting them as isolated techniques. Learning about preprocessing, statistical representations, and semantic language models helped me see that NLP is a structured process where each stage contributes to improving the quality of the final output. Instead of memorizing individual concepts, I began to understand how they connect and why they are necessary. This shift in perspective made the topic feel much more practical and relevant to my studies as an Information Technology student.

II. EXPERIENTIAL AND PRACTICAL INSIGHTS

The hands-on text analytics exercises gave me a clearer understanding of concepts that initially seemed abstract in the readings. Working with textual data demonstrated that even small preprocessing decisions, such as removing stop

words or applying lemmatization, can influence the quality of the analysis. Seeing these effects through practical examples made the learning process much more meaningful than simply reading about the techniques.

Another valuable insight was recognizing that there is no single approach that works best for every NLP problem. Some methods are computationally simple and easy to interpret, while others provide richer contextual understanding at the cost of increased complexity. This experience encouraged me to think more critically about selecting appropriate techniques based on the nature of the data and the objectives of a project rather than assuming that newer methods are always the best solution.

III. CORE LEARNING OUTCOMES

One of my biggest takeaways from this unit was understanding the trade-offs involved in tokenization and text preprocessing. Although preprocessing improves consistency, I learned that excessive cleaning may also remove information that contributes to the meaning of a sentence. This made me realize that preprocessing decisions should always be aligned with the intended application rather than applied mechanically.

I also developed a much clearer understanding of the differences between TF-IDF and word embeddings. TF-IDF is effective for identifying terms that are important within documents because it measures word importance statistically. However, it does not capture the relationships or meanings between words. In contrast, embeddings represent words as vectors within a semantic space, allowing models to recognize similarities and contextual relationships that traditional statistical methods cannot represent.

Perhaps the most interesting concept for me was context vectors. Learning that modern language models generate different representations for the same word depending on its surrounding context completely changed my perspective on how AI understands language. Instead of treating words as isolated units, contextual models interpret meaning based on the entire sentence, making modern NLP systems significantly more flexible and accurate.

IV. COURSE APPLICATION AND IMPACT

This unit has given me a stronger foundation for the remaining parts of the course. Understanding the preprocessing pipeline and different text representation methods will make it easier to work with future laboratory activities involving machine learning and more advanced NLP techniques. Rather than treating algorithms as "black boxes," I now have a better understanding of the processes that occur before a model generates predictions.

I also see direct connections between these concepts and my capstone project. Although my current project focuses on an emergency response and resource management system, many of the ideas introduced in this unit can be applied to analyzing incident reports, extracting useful information from textual records, and supporting decision-making through intelligent text processing. Learning these fundamentals gives me greater confidence that I can integrate AI techniques into future system development when appropriate.

V. CONCLUSION AND FUTURE WORK

Overall, Unit 5 significantly expanded my understanding of Natural Language Processing. More than learning individual

techniques, I gained a better appreciation of how preprocessing, statistical methods, and semantic representations work together to transform raw text into meaningful information. The practical exercises reinforced the theoretical concepts and showed me that effective NLP depends not only on sophisticated models but also on thoughtful preparation and careful representation of data.

Moving forward, I would like to explore more advanced topics such as transformer architectures, attention mechanisms, Retrieval-Augmented Generation (RAG), and large language models. These technologies are becoming increasingly important across many areas of artificial intelligence, and I believe that building on the strong foundation established in this unit will help me understand and apply them more effectively in future coursework, research, and software development projects.

REFERENCES

[Introduction to natural language processing concepts - Training | Microsoft Learn](#)